


Figure 12. The definitive test: on a clean, matched-pipeline, larger sample the "Arabic compression" direction does NOT hold — Arabic is significantly WIDER on the cross-encoder/MiniLM and indistinguishable on LaBSE. The geometry gap is real but its direction is unstable (corpus-, genre-, and model-dependent) — which is exactly why the cutoff must be calibrated per tenant.

(a) The direction flips with corpus size/diversity



(b) Clean matched-pipeline (n=390): no compression

